

(S8586)

Writing Graph Primitives with Gunrock

Muhammad Osama & Yuechao Pan
w/ The Gunrock Team



<https://gunrock.github.io>

UC DAVIS
UNIVERSITY OF CALIFORNIA

Why Graph Processing?

ted:

Twitter Dataset 1 Overview

ds

Tweets: 292.7 Million +
Unique Users: 7,619,916
Total Size: 232 GB

are

s.

API
et
n,

Figure 1: Collection profile of Twitter Dataset 1

Twitter Dataset 2 Overview

Tweets: 1 Billion+
Unique Users: 94 Million+
Geolocated Tweets: 31 Million+
Total Size: 146 GB



Figure 2: Twitter in Europe
Image obtained from <https://blog.twitter.com/2013/geography-tweets-3>



Figure 1: IPv4 Census Map (<http://www.caida.org/research/id-consumption/census-map/images/20061108.png>)

Data Type	# of Records (Size)	Quick Description
Service Probes	180 billion (5.5 TB)	Results of probes with different formats sent to various service ports of IPv4 addresses.
Reverse DNS	10.5 billion (366 GB)	Results of DNS name requests (reverse lookups) for addresses within the IPv4 space using 16 large DNS Servers.
TCP/IP Fingerprints	80 million (50 GB)	Results of remote OS detection fingerprinting from NMap tool.

Table 1: Net Data Volume



Figure 1: Map of Jobs (Colored by Country)

Background and Formats: The dataset consists of 119+ Million jobs and is about 40 GB in size. There are approximately 2.1 million unique jobs in the set as many records are duplicates. To

Data Field	Example
Posted Date	2012-10-23
Location	Capital Federal
Department	Capital Federal

Bitcoin Data Set Overview (May 15, 2013)

Transactions: 15.8 Million+
Edges: 37.4 Million +
Senders: 5.4 Million+
Receivers: 6.3 Million+
Bitcoins Transacted: 1.4 Million +

Why use GPUs for Graph Processing?

Graphs

- Found everywhere
 - Road & social networks, web, etc.
- Require fast processing
 - Memory bandwidth, computing power and GOOD software
- Becoming very large
 - Billions of edges
- Irregular data access pattern and control flow
 - Limits performance and scalability

GPUs

- Found everywhere
 - Data center, desktops, mobiles, etc.
- Very powerful
 - High memory bandwidth (900 GBps) and computing power (15.7 TFlops)
- Limited memory size
 - 16 GB per NVIDIA V100
- Hard to program
 - Harder to optimize

What is Gunrock?

What is the Gunrock Library?

A CUDA-based graph processing library, aims for:

- **Generality**
covers a broad range of graph algorithms
- **Performance**
maintains good performance
- **Programmability**
Makes it easy to implement graph algorithms
extends from 1-GPU to multi-GPUs as simple as possible
- **Scalability**
fits in (very) limited GPU memory space
performance scales when using more GPUs

gunrock / gunrock

Unwatch 62 Unstar 308 Fork 86

Code Issues 46 Pull requests 0 Projects 0 Wiki Insights Settings

High-Performance Graph Primitives on GPUs Edit

gunrock cuda graph-processing graph-analytics gpu graph-primitives Manage topics

1,538 commits 5 branches 5 releases 18 contributors Apache-2.0

Branch: dev-refactor New pull request Create new file Upload files Find file Clone or download

This branch is 151 commits ahead, 56 commits behind master. Pull request Compare

sgpyc Merge pull request #330 from sgpyc/dev-refactor Latest commit 86a8c16 17 hours ago

cmake	[util::Parameters] moved fully specialized code into types.cu	a year ago
codesnaps	[SSSP] added explanation to different parts of the mGPU framework	a year ago
dataset	Fixed *_sync() bug for CUDA 9, it now takes mask as a parameter as well.	6 months ago
dep	Updated installation to boost 1.58.0	2 years ago
doc	rapidjson, release notes and porting guide (#306)	2 months ago
externals	Dev refactor (#305)	2 months ago
gunrock	mached types	17 hours ago
output	looks like i added dobfs/bfs support	2 years ago

<https://gunrock.github.io>

How does Gunrock work?

Gunrock: Programming Model


Data-centric abstraction

- A **frontier**; group of vertices or edges
- Manipulation of frontiers is an **operation**

Bulk-synchronous programming


- Series of **parallel operations** separated by **global barriers**

Gunrock: Programming Model (cont.)




Advance

Generates new frontier by visiting the neighbors.



Filter


Chooses a **subset** of current frontier as the new front.



Compute

Applies a **compute** operation on all elements in its input front.

Gunrock: Programming Model (cont.)




How do you write a
graph primitive using Gunrock?

Example:

The Shortest Path Problem


Find a path between two vertices in a graph such that the sum of the weights of its constituent edges is minimized.



Example:

The Shortest Path Problem

- Let **A** be the source node.
- Shortest Path: (A, **C**, **E**, **D**, **F**)



Example:

Single Source Shortest Path

In SSSP we have to find **shortest paths** from a **source vertex v** to **all other vertices** in the graph.

Example: Gunrock's SSSP Algorithm Pseudocode

```
1: function SSSP(Graph, source):
2:   for each vertex v in Graph:           // Initialization
3:     distances[v] := infinity; predecessors[v] := undefined
4:   distances[source] := 0                 // Distance of source
5:   Q0 := {source}                       // Set of vertices (frontier)
6:
7:   while Q0 is not empty:               // Iteration loop
8:     Q1 := {}
9:     for each v in Q0, do in parallel:
10:      for each edge e<v,u> of v, do in parallel:
11:        new_distance := distances[v] + weights[e]
12:        temp_distance := min(distances[u], new_distance) // atomic
13:        if (temp_distance < distances[u])
14:          predecessors[u] := v
15:          put u in Q1
16:        else
17:          put invalid in Q1
18:     Q0 := Q1
19:     for each u in Q1, do in parallel:
20:       if (u is valid):
21:         put u in Q0
```

Reset { 2: for each vertex v in Graph: // Initialization
3: distances[v] := infinity; predecessors[v] := undefined
4: distances[source] := 0 // Distance of source
5: Q₀ := {source} // Set of vertices (frontier)

Advance { 7: while Q₀ is not empty: // Iteration loop
8: Q₁ := {}
9: for each v in Q₀, do in parallel:
10: for each edge e<v,u> of v, do in parallel:
11: new_distance := distances[v] + weights[e]
12: temp_distance := min(distances[u], new_distance) // atomic
13: if (temp_distance < distances[u])
14: predecessors[u] := v
15: put u in Q₁
16: else
17: put invalid in Q₁
18: Q₀ := Q₁

Filter { 19: for each u in Q₁, do in parallel:
20: if (u is valid):
21: put u in Q₀

Directory Structure

gunrock/...

Gunrock (Repository) at <https://gunrock.github.io>

- **gunrock**
 - app
 - **<primitive>**
 - <primitive>_problem.cuh
 - <primitive>_enactor.cuh
 - <primitive>_test.cuh
 - <primitive>_app.cu
 - ...
- **tests**
 - <primitive>
 - test_<primitive>.cu
 - ...
- ...



gunrock/app/<primitive> directory

Our new gunrock primitive will be defined in this directory, along with the tests associated with it.



gunrock/tests/<primitive> directory

A test driver file to create an executable for your primitive.

Directory Structure

gunrock/app/<primitive>

Gunrock (Repository) at <https://gunrock.github.io>

- **gunrock**
 - app
 - <primitive>
 - <primitive>_problem.cuh → GPU storage management structure for primitive's problem data.
 - <primitive>_enactor.cuh → Enactor file that operates on the data.
 - <primitive>_test.cuh → CPU implementation for comparison or correctness checking.
 - <primitive>_app.cu → Higher level routines for the primitive.

Example: SSSP Problem


```
1: function SSSP(Graph, source):  
2:   for each vertex v in Graph:  
3:     distances[v] := infinity; predecessors[v] := undefined  
4:   distances[source] := 0  
5:   Q0 := {source}  
6:  
...
```

Reset()


BaseDataSlice()

Example: SSSP Enactor

```
1: function SSSP(Graph, source):  
  ...  
7:   while  $Q_0$  is not empty:  
8:      $Q_1 := \{\}$   
9:     for each  $v$  in  $Q_0$ , do in parallel:  
10:      for each edge  $e<v,u>$  of  $v$ , do in parallel:  
11:        new_distance := distances[ $v$ ] + weights[ $e$ ]  
12:        temp_distance := min(distances[ $u$ ], new_distance)  
13:        if (temp_distance > new_distance)  
14:          predecessors[ $u$ ] :=  $v$   
15:          put  $u$  in  $Q_1$   
16:        else  
17:          put invalid in  $Q_1$   
18:      $Q_0 := \{\}$   
19:     for each  $u$  in  $Q_1$ , do in parallel:  
20:       if ( $u$  is valid):  
21:         put  $u$  in  $Q_0$ ;
```



Advance



Filter

Example:

SSSP Test

- Display solution
- Result validation
 - CPU Reference
 - Correctness checking

Example:

SSSP App

- Call to **Problem** and **Enactor**
- Recording Statistics
- External APIs

Interface to C

Interface to Python

Coding Tutorial

available @ <https://tinyurl.com/gunrockgtc>

Section S8594, Wednesday, Mar 28, 3:00–3:25 PM, Hilton Santa Clara

Latest Development of Gunrock: a Graph Processing Library on GPUs

- New APIs
 - Operators: Advance, filter, compute
 - External interfaces
 - Graph representations
- New Graph Applications
 - Subgraph matching
 - Graph coloring
 - Random walks, etc.

Acknowledgements

The Gunrock team

Gunrock code contributors

NVIDIA

For hardware support, GPU cluster access, and all other support and discussion

Funding support:

- DARPA HIVE program
- DARPA XDATA program under US Army award W911QX-12-C-0059
- DARPA STTR awards D14PC00023 and D15PC00010
- NSF awards OAC-1740333 and CCF-1629657
- Adobe Data Science Research Award

Questions?

Q: How can I find Gunrock?

A: <https://gunrock.github.io/>

Q: Is it free and open?

A: Absolutely (under Apache License v2.0)

Q: Papers, slides, etc.?

A: <http://gunrock.github.io/gunrock/doc/latest/index.html#Publications>

Q: Requirements?

A: **CUDA \geq 8.0**, GPU compute capability \geq 3.0, Linux || Mac OS

Q: Language?

A: C/C++, with a simple wrapper to connect to Python